

Project 2

CS 5/7394 - Applied Machine Learning

Sofia Murillo, Randolph Rankin

- **Due** - March 11 @ 11:59 pm pushed to Github repo
- **Teams** - You can do this project solo or in pairs. Not 3, not 4 not 5... Max of 2. If a 5394 student pairs with a 7394 student, the pair needs to do the 7394 work.

Below are 6 Kaggle Datasets. You will choose 1 to work with for this project.

- [Airfare Prediction Dataset](#)
- [Chinese Rest Holiday Dataset](#)
- [Jigsaw Toxic Comment Classification Challenge](#)
- [Latest Covid 19 Dataset Worldwide](#)
- [Trains](#)
- [Football Data top 5 Leagues](#) (*Selected Dataset*)

Merging disparate datasets is a staple of the data exploration process. Therefore, for which ever data set above that you choose, you will need to independently find **an additional** dataset to merge with your selection. The only requirement is that it add to the richness of the original dataset. Students in the 7000-level version of the class need to find two additional data sets to merge with the original selection.

Note: If you want to start with a different data set, you need to get Fontenot's OK first.

Your Tasks

Below, there are cells that provide directions on what to do for the project.

You can insert as many cells between the ones below as you'd like, but please **Do NOT** change the cells already provided.

Part 1 - Getting Started

- Import libraries
- Load original Data (which ever one you chose from the provided list) into a data frame.
- Load your additional data set(s) into a data frame.
- In a markdown cell, provide a brief description of your the data sets you've chosen to work with.
- Develop a list of 3 - 4 questions that you hope to be able to answer after the exploration of the data and write them in this section.

```
In [1]: # "Import libraries"

import pandas as pd
import numpy as np
import math
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib.style as style

import warnings
warnings.filterwarnings('ignore')
```

```
style.use('seaborn')
```

```
In [2]: # "Load original Data (which ever one you chose from the provided list) into a data frame.
original_data = pd.read_csv('data/combined_data.csv')

# "Load your additional data set(s) into a data frame."
additional_data = pd.read_excel('data/all-leaguetables.xlsx')
```

Data Explanation

Primary Dataset

Our primary dataset consists of match-level data for 5 top soccer (football) leagues. Data for each match is related to scores and specific game statistics relevant to each match such as ball possession and penalties. This dataset provides insight to overall match statistics.

```
In [3]: original_data.head()
```

Out[3]:

	Unnamed: 0	Home Team	Away Team	Score	Half Time Score	Match Excitement	Home Team Rating	Away Team Rating	Home Team Possession %	Away Team Possession %
0	0	MAN UTD	SWANSEA	1-2	0-1	5.9	5.6	7.6	60	40
1	1	WEST BROM	SUNDERLAND	2-2	1-1	7.3	6.5	7.4	58	42
2	2	LEICESTER CITY	EVERTON	2-2	1-2	7.0	6.5	6.3	37	63
3	3	WEST HAM	TOTTENHAM	0-1	0-0	4.8	5.9	6.4	47	53
4	4	QPR	HULL CITY	0-1	0-0	3.8	5.7	6.6	51	49

5 rows x 42 columns

Additional Dataset

Our additional dataset contains data pertaining to season-level record data for several football leagues. This data itself contains the statistics related to a team's performance over a given season, including total wins/draws/losses, league rankings, and cumulative goals scored. The dataset provides insight into the overall performance of each team in a given season.

```
In [4]: additional_data.head()
```

Out[4]:

	League	Season	Position	Team	Matches	Wins	Draws	Losses	GoalsScored	GoalsConceded	Gc
0	Premier League	2015/16	2	Arsenal	38	20	11	7	65		36
1	Premier League	2015/16	20	Aston Villa	38	3	8	27	27		76
2	Premier League	2015/16	16	Bournemouth	38	11	9	18	45		67

	League	Season	Position	Team	Matches	Wins	Draws	Losses	GoalsScored	GoalsConceded	Gc
3	Premier League	2015/16	10	Chelsea	38	12	14	12	59	53	
4	Premier League	2015/16	15	Crystal Palace	38	11	9	18	39	51	

Questions Guiding Our Exploration

1. Are match statistics good indicators of a team's performance in the future?
2. Does a correlation exist between a team's season ranking and certain match statistics?
3. Are match wins truly influenced by stats, or are there external conditions that cannot be found within the data?
4. Can the relative "excitement" of a match along with a home team advantage predict a team's performance?

Part 2 - Data Inspection

Write some code to summarize the datasets. Think about the following questions:

- What type of data is each variable? (think like a data scientist here, not a computer scientist)
- What is the total size of the data sets?
- What time boundaries are there in the dataset? IOW, what time frame do they span?
- Are there any missing values in any of the variables?

Do this with Intentionality. Don't skimp.

Describing the Original Dataset

By previewing our dataframe and calling `DataFrame.info()`, we can see each variable's datatype as well as any missing variables, along with gaining an understanding of the size of our data.

In [5]: `original_data.head()`

Out[5]:

	Unnamed: 0	Home Team	Away Team	Score	Half Time Score	Match Excitement	Home Team Rating	Away Team Rating	Home Team Possession %	Away Team Possession %
0	0	MAN UTD	SWANSEA	1-2	0-1	5.9	5.6	7.6	60	40
1	1	WEST BROM	SUNDERLAND	2-2	1-1	7.3	6.5	7.4	58	42
2	2	LEICESTER CITY	EVERTON	2-2	1-2	7.0	6.5	6.3	37	63
3	3	WEST HAM	TOTTENHAM	0-1	0-0	4.8	5.9	6.4	47	53
4	4	QPR	HULL CITY	0-1	0-0	3.8	5.7	6.6	51	49

5 rows × 42 columns

In [6]: `original_data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12062 entries, 0 to 12061
Data columns (total 42 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	12062 non-null	int64
1	Home Team	12062 non-null	object
2	Away Team	12062 non-null	object
3	Score	12062 non-null	object
4	Half Time Score	12062 non-null	object
5	Match Excitement	12062 non-null	float64
6	Home Team Rating	12062 non-null	float64
7	Away Team Rating	12062 non-null	float64
8	Home Team Possession %	12062 non-null	int64
9	Away Team Possession %	12062 non-null	int64
10	Home Team Off Target Shots	12062 non-null	float64
11	Home Team On Target Shots	12062 non-null	float64
12	Home Team Total Shots	12062 non-null	float64
13	Home Team Blocked Shots	12062 non-null	float64
14	Home Team Corners	12062 non-null	float64
15	Home Team Throw Ins	12062 non-null	float64
16	Home Team Pass Success %	12062 non-null	float64
17	Home Team Aerials Won	12062 non-null	float64
18	Home Team Clearances	12062 non-null	float64
19	Home Team Fouls	12062 non-null	float64
20	Home Team Yellow Cards	12062 non-null	float64
21	Home Team Second Yellow Cards	12062 non-null	float64
22	Home Team Red Cards	12062 non-null	float64
23	Away Team Off Target Shots	12062 non-null	float64
24	Away Team On Target Shots	12062 non-null	float64
25	Away Team Total Shots	12062 non-null	float64
26	Away Team Blocked Shots	12062 non-null	float64
27	Away Team Corners	12062 non-null	float64
28	Away Team Throw Ins	12062 non-null	float64
29	Away Team Pass Success %	12062 non-null	float64
30	Away Team Aerials Won	12062 non-null	float64
31	Away Team Clearances	12062 non-null	float64
32	Away Team Fouls	12062 non-null	float64
33	Away Team Yellow Cards	12062 non-null	float64
34	Away Team Second Yellow Cards	12062 non-null	float64
35	Away Team Red Cards	12062 non-null	float64
36	Home Team Goals Scored	12062 non-null	int64
37	Away Team Goals Scored	12062 non-null	int64
38	Home Team Goals Conceded	12062 non-null	int64
39	Away Team Goals Conceded	12062 non-null	int64
40	year	12062 non-null	int64
41	league	12062 non-null	object

```
dtypes: float64(29), int64(8), object(5)
```

```
memory usage: 3.9+ MB
```

Data is accounted for in each row of our dataframe, suggesting no missing values.

```
In [7]: original_data['year'].unique()
```

```
Out[7]: array([2014, 2015, 2016, 2017, 2018, 2019, 2020])
```

The timeframe for our match-level data spans from 2014-2020

```
In [8]: # Obtain the number of teams and home matches contained in our primary dataset
original_data['Home Team'].value_counts()
```

```
Out[8]: MAN UTD      133
SAMPDORIA    133
MAN CITY     133
```

```

VALENCIA      133
FIORENTINA    133
...
NÜRNBERG      17
DARMSTADT     17
PADERBORN     17
BIELEFELD     17
INGOLSTADT    17
Name: Home Team, Length: 146, dtype: int64

```

From here, we can see that there is a total of 146 teams within the original dataset, each of which playing between 17 and 133 home matches over the 7 years of data.

```

In [9]: # Check that the teams contained in the "away" column match the teams in the "home" column
not False in original_data['Away Team'].sort_values().unique() == original_data['Home Team'].sort_values().unique()

```

```

Out[9]: True

```

By checking that the unique lists of both Home and Away team columns are equivalent, we can confirm that all 146 teams are accounted for as both home and away teams in this dataset. This allows us the assumption that a unique list of either column is the complete list of all teams.

Describing the Additional Dataset

Again, we'll preview our data and use `DataFrame.info()` to explore the variables contained in our season-level record data.

```

In [10]: additional_data.head()

```

```

Out[10]:

```

	League	Season	Position	Team	Matches	Wins	Draws	Losses	GoalsScored	GoalsConceded	Gc
0	Premier League	2015/16	2	Arsenal	38	20	11	7	65	36	
1	Premier League	2015/16	20	Aston Villa	38	3	8	27	27	76	
2	Premier League	2015/16	16	Bournemouth	38	11	9	18	45	67	
3	Premier League	2015/16	10	Chelsea	38	12	14	12	59	53	
4	Premier League	2015/16	15	Crystal Palace	38	11	9	18	39	51	

```

In [11]: additional_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 399 entries, 0 to 398
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   League          399 non-null   object
1   Season          399 non-null   object
2   Position        399 non-null   object
3   Team            399 non-null   object
4   Matches         399 non-null   int64
5   Wins            399 non-null   int64
6   Draws           399 non-null   int64
7   Losses          399 non-null   int64

```

```

8   GoalsScored      399 non-null    int64
9   GoalsConceded    399 non-null    int64
10  GoalDiff          399 non-null    int64
11  FinalPoints       399 non-null    int64
dtypes: int64(8), object(4)
memory usage: 37.5+ KB

```

As with the primary dataset, no values appear to be missing from any row.

```
In [12]: additional_data['Season'].unique()
```

```
Out[12]: array(['2015/16', '2016/17', '2017/18', '2018/19', '2019/20', 2015, 2016,
        2017, 2018, 2019], dtype=object)
```

Data in our season-level data ranges from 2015 to 2019.

```
In [13]: additional_data['Team'].value_counts()
```

```
Out[13]: Arsenal                5
Bayer Leverkusen              5
Chicago Fire                  5
Real Madrid                   5
Real Sociedad                 5
..
FC Union Berlin               1
SD Huesca                     1
RCD Mallorca                  1
FC Nurnberg                   1
FC Cincinnati                 1
Name: Team, Length: 104, dtype: int64
```

Exploring this second dataset, we see a different number of teams reported. This informs us that league data will need to be reconciled between the two datasets during our cleaning process.

Part 3 - Data Description

- Create a data description (data dictionary) for your data sets.
 - Describe each variable
 - If categorical, what levels are present? If the levels are encoded, what do the codes mean?
 - If numeric, provide min, max, median and any other univariate stats you'd like to add in.
- Where appropriate, provide histograms or other visualizations to characterize each variable.

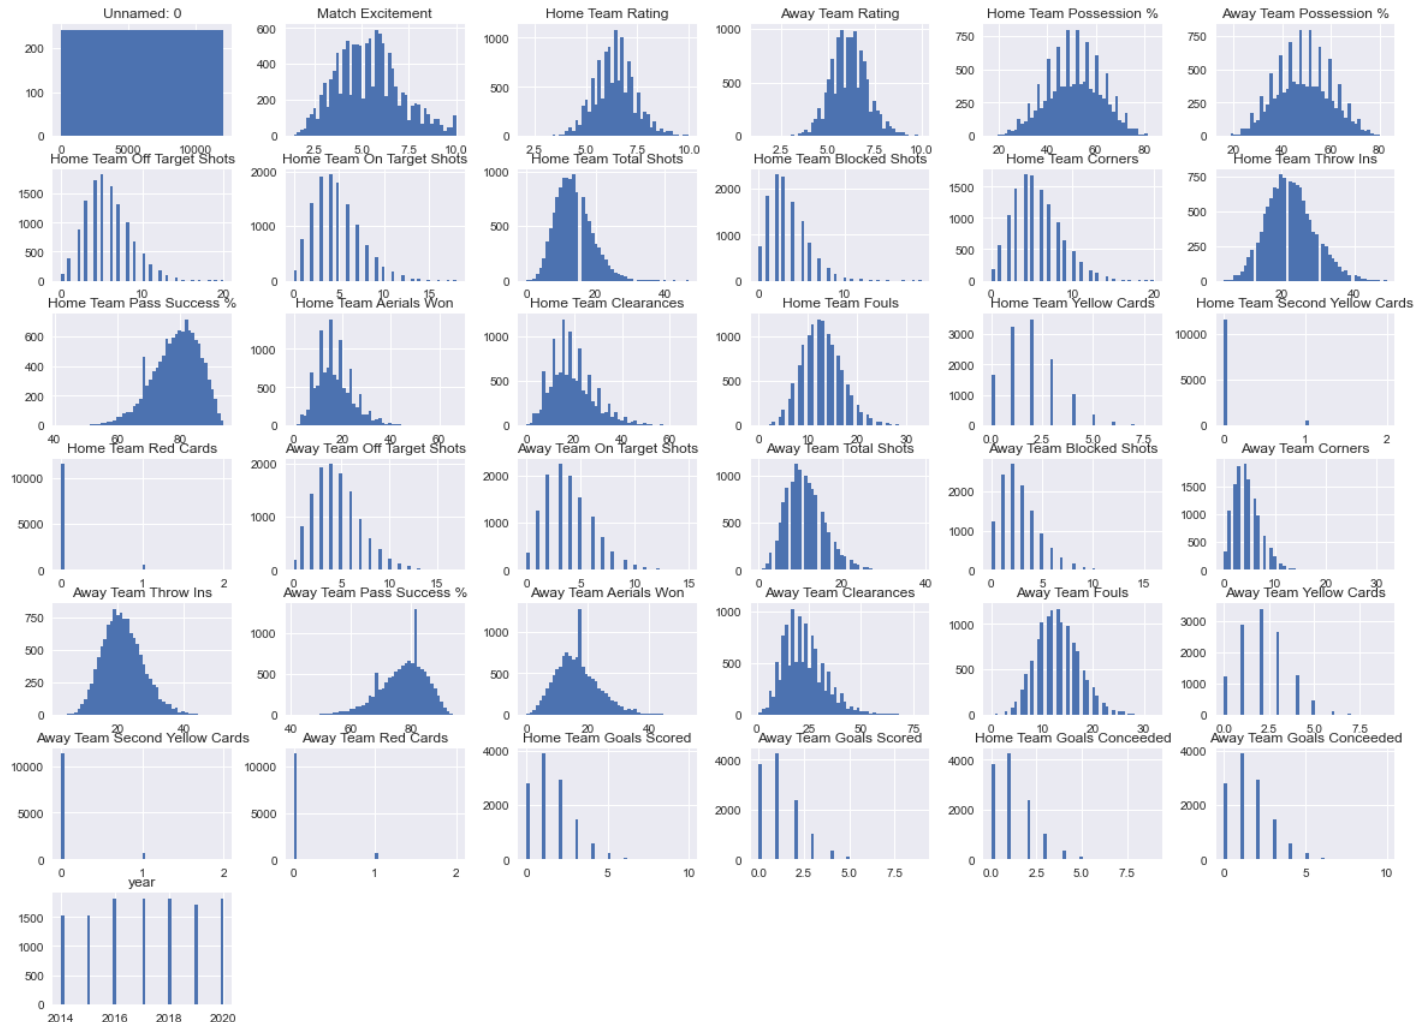
Data Dictionary - Original Data

Variable	Type	Description
Home Team	Categorical (Nominal)	Name of the hosting team
Away Team	Categorical (Nominal)	Name of the visiting team
Score	Numerical (Discrete, Compound Variable)	Final score of the match, delimited by a hyphen, with the home team score on the left side, and the away team score on the right
Half Time Score	Numerical (Discrete, Compound Variable)	Similar to the score variable, but representing the score after the first half
Match Excitement	Numerical (Continuous)	Index calculated between 0 and 10 based on game statistics that would indicate an "exciting" match
Home / Away Team Rating	Numerical (Continuous)	Index representing relative performance of a team prior to a match
Home / Away Team Possession %	Numerical (Discrete*)	Percentage of total gameplay time for which the team had control of the ball

Variable	Type	Description
Home / Away Team Off Target Shots	Numerical (Discrete)	Total off target shots taken by a team
Home / Away Team On Target Shots	Numerical (Discrete)	Total on target shots taken by a team
Home / Away Team Total Shots	Numerical (Discrete)	Total shots attempted by a team
Home / Away Team Blocked Shots	Numerical (Discrete)	Total shots attempted by a team that were blocked
Home / Away Team Corners	Numerical (Discrete)	Total corner kicks attempted by a team
Home / Away Team Throw Ins	Numerical (Discrete)	Total number of throw-ins initiated by a team
Home / Away Team Pass Success %	Numerical (Discrete*)	Percentage of successful passes out of total attempted passes
Home / Away Team Aerials Won	Numerical (Discrete)	Number of successful aerial passes completed
Home / Away Team Clearances	Numerical (Discrete)	Number of times a team cleared the ball away from their goal defensively
Home / Away Team Fouls	Numerical (Discrete)	Number of fouls committed by a team
Home / Away Team Yellow Cards	Numerical (Discrete)	Number of yellow cards issued to a team's players
Home / Away Team Second Yellow Cards	Numerical (Discrete)	Number of second yellow cards issued to a team's player
Home / Away Team Red Cards	Numerical (Discrete)	Number of red cards issued a teams players
Home / Away Team Goals Scored	Numerical (Discrete)	Number of successful goals scored by a team
Home / Away Team Goals Conceded (sic)	Numerical (Discrete)	Number of failed blocks of the opposing team's goal attempts
year	Categorical (Ordinal)	Season year during which the match occurred
league	Categorical (Nominal)	League of which the playing teams are members

* *Would otherwise be continuous, but the data is presented rounded to the nearest whole percent*

```
In [14]: original_data.hist(bins=50, figsize=(20,15))
pass
```




Data Dictionary - Additional Data

Variable	Type	Description
League	Categorical (Nominal)	League of which the playing teams are members
Season	Categorical (Ordinal)	Season year for which the statistics apply
Position	Categorical (Ordinal)	Rank of the team amongst all teams in their league
Team	Categorical (Nominal)	Name of the team for which the statistics apply
Matches	Numerical (Discrete)	Number of matches played by the team for this season
Wins	Numerical (Discrete)	Number of matches won by the team for this season
Draws	Numerical (Discrete)	Number of matches played by the team that ended in draws for this season
Losses	Numerical (Discrete)	Number of matches lost by the team for this season
GoalsScored	Numerical (Discrete)	Number of total goals scored by the team for this season
GoalsConceded	Numerical (Discrete)	Number of total goals scored against the team for this season
GoalDiff	Numerical (Discrete)	Difference between GoalsScored and GoalsConceded

Variable	Type	Description
FinalPoints	Numerical (Discrete)	Number of aggregate points as determined by league rules regarding multiple matches between the same teams

```
In [15]: additional_data.hist(bins=50, figsize=(20,15))
pass
```



Part 4 - Merge the data

Now that you have a better feel for each of your two (or three, for the 7394 students) data sets, it is time to merge them. Describe your strategy for merging the data sets and then actually perform the merge.

Develop a strategy for verifying that the data is properly merged (hoping and finger-crossing are not valid strategies).

Our Merging Strategy

When cleaning the data for merging, we'll take a couple of primary variables into account—the season of the match and the teams who played. We'll first isolate only the leagues for which data is common between the two sets. Then, we'll explore how we'll clean the string variables for the team names to be consistent with one another.

To verify our merge is successful, we'll calculate season records based on match wins and losses from our original data set and test that the records as calculated on a match level are consistent with the season

records as reported in our additional data set. This will ensure that each merged row is consistent with the data from both sets.

League Filtering

First, we'll want to explore which leagues are reported in each of our data.

```
In [16]: original_data.league.unique()

Out[16]: array(['english', 'french', 'german', 'italian', 'spanish'], dtype=object)
```

```
In [17]: additional_data.League.unique()

Out[17]: array(['Premier League', 'La Liga', 'Bundesliga', 'Major League Soccer'],
              dtype=object)
```

We see that in our additional data, we have unneeded data for Major League Soccer, while lacking data for Italian and French leagues.

```
In [18]: df_orig = original_data.loc[original_data.league.isin(['english', 'german', 'spanish'])]
df_addition = additional_data.loc[additional_data.League.isin(['Premier League', 'La Liga'])]
```

Team Name Formatting

Next, we'll explore how the team names are formatted differently for data cleaning purposes.

```
In [19]: df_addition.Team = df_addition.Team.apply(lambda x: str.upper(x))
```

Now, let's see which teams can immediately be merged...

```
In [20]: # Values that are matched between datasets
df_orig['Home Team'].loc[df_orig['Home Team'].isin(df_addition.Team)].sort_values().unique()

Out[20]: array(['ARSENAL', 'ASTON VILLA', 'ATLETICO MADRID', 'BOURNEMOUTH',
              'BRIGHTON', 'BURNLEY', 'CHELSEA', 'CRYSTAL PALACE', 'ESPANYOL',
              'EVERTON', 'FULHAM', 'HULL CITY', 'LEICESTER CITY', 'LIVERPOOL',
              'MIDDLESBROUGH', 'RAYO VALLECANO', 'RB LEIPZIG', 'REAL MADRID',
              'REAL SOCIEDAD', 'SEVILLA FC', 'SOUTHAMPTON', 'SUNDERLAND',
              'WATFORD'], dtype=object)
```

... and which teams will need reformatting.

```
In [21]: # Values that are unmatched in our additional data
df_addition.Team.loc[~df_addition.Team.isin(df_orig['Home Team'])].sort_values().unique()

Out[21]: array(['ATHLETIC BILBAO', 'BAYER LEVERKUSEN', 'BAYERN MUNICH',
              'BORUSSIA DORTMUND', 'BORUSSIA MONCHENGLADBACH', 'CA OSASUNA',
              'CARDIFF CITY', 'CD LEGANES', 'CELTA VIGO', 'DEPORTIVO ALAVES',
              'DEPORTIVO LA CORUNA', 'EINTRACHT FRANKFURT', 'FC AUGSBURG',
              'FC BARCELONA', 'FC INGOLSTADT 04', 'FC KOLN', 'FC NURNBERG',
              'FC SCHALKE 04', 'FC UNION BERLIN', 'FORTUNA DUSSELDORF',
              'FSV MAINZ 05', 'GETAFE CF', 'GIRONA FC', 'GRANADA CF',
              'HAMBURGER SV', 'HANNOVER 96', 'HERTHA BSC', 'HUDDERSFIELD TOWN',
              'LEVANTE UD', 'MALAGA CF', 'MANCHESTER CITY', 'MANCHESTER UNITED',
              'NEWCASTLE UNITED', 'NORWICH CITY', 'RCD MALLORCA',
              'REAL BETIS BALOMPIE', 'REAL VALLADOLID CF', 'SC FREIBURG',
              'SC PADERBORN 07', 'SD EIBAR', 'SD HUESCA', 'SHEFFIELD UNITED',
```

```
'SPORTING GIJON', 'STOKE CITY', 'SV DARMSTADT 98',
'SV WERDER BREMEN', 'SWANSEA CITY', 'TOTTENHAM HOTSPUR',
'TSG 1899 HOFFENHEIM', 'UD LAS PALMAS', 'VALENCIA CF',
'VFB STUTTGART', 'VFL WOLFSBURG', 'VILLARREAL CF',
'WEST BROMWICH ALBION', 'WEST HAM UNITED',
'WOLVERHAMPTON WANDERERS']], dtype=object)
```

```
In [22]: # Values that are unmatched in our original data
df_orig['Home Team'].loc[~df_orig['Home Team'].isin(df_addition.Team)].sort_values().unique()
```

```
Out[22]: array(['1. FC KÖLN', 'ALAVÉS', 'ALMERÍA', 'ATHLETIC', 'AUGSBURG',
'BARCELONA', 'BAYERN', 'BIELEFELD', 'CARDIFF', 'CELTA', 'CÁDIZ CF',
'CÓRDOBA', 'DARMSTADT', 'DEPORTIVO', 'DORTMUND', 'DÜSSELDORF',
'EIBAR', 'ELCHE', 'FRANKFURT', 'FREIBURG', 'GETAFE', 'GIJÓN',
'GIRONA', 'GRANADA', 'HANNOVER', 'HERTHA', 'HOFFENHEIM', 'HSV',
'HUDDERSFIELD', 'HUESCA', 'INGOLSTADT', 'LAS PALMAS', 'LEEDS UTD',
'LEGANÉS', 'LEVANTE', 'LEVERKUSEN', 'M'GLADBACH', 'MAINZ',
'MALLORCA', 'MAN CITY', 'MAN UTD', 'MÁLAGA', 'NEWCASTLE',
'NORWICH', 'NÜRNBERG', 'OSASUNA', 'PADERBORN', 'QPR', 'REAL BETIS',
'SCHALKE', 'SHEFFIELD UTD', 'STOKE', 'STUTTGART', 'SWANSEA',
'TOTTENHAM', 'UNION BERLIN', 'VALENCIA', 'VALLADOLID',
'VILLARREAL', 'W. BREMEN', 'WEST BROM', 'WEST HAM', 'WOLFSBURG',
'WOLVES'], dtype=object)
```

For our mismatched team names, we'll replace the names in our additional dataset to match the formatting of our base dataset.

```
In [23]: # Dictionary mapping current team name in additional data set to corresponding name in original
team_name_key = {
    'ATHLETIC BILBAO': 'ATHLETIC',
    'BAYER LEVERKUSEN': 'LEVERKUSEN',
    'BAYERN MUNICH': 'BAYERN',
    'BORUSSIA DORTMUND': 'DORTMUND',
    'BORUSSIA MONCHENGLADBACH': 'M'GLADBACH',
    'CA OSASUNA': 'OSASUNA',
    'CARDIFF CITY': 'CARDIFF',
    'CD LEGANES': 'LEGANÉS',
    'CELTA VIGO': 'CELTA',
    'DEPORTIVO ALAVES': 'ALAVÉS',
    'DEPORTIVO LA CORUNA': 'DEPORTIVO',
    'EINTRACHT FRANKFURT': 'FRANKFURT',
    'FC AUGSBURG': 'AUGSBURG',
    'FC BARCELONA': 'BARCELONA',
    'FC INGOLSTADT 04': 'INGOLSTADT',
    'FC KOLN': '1. FC KÖLN',
    'FC NURNBERG': 'NÜRNBERG',
    'FC SCHALKE 04': 'SCHALKE',
    'FC UNION BERLIN': 'UNION BERLIN',
    'FORTUNA DUSSELDORF': 'DÜSSELDORF',
    'FSV MAINZ 05': 'MAINZ',
    'GETAFE CF': 'GETAFE',
    'GIRONA FC': 'GIRONA',
    'GRANADA CF': 'GRANADA',
    'HAMBURGER SV': 'HSV',
    'HANNOVER 96': 'HANNOVER',
    'HERTHA BSC': 'HERTHA',
    'HUDDERSFIELD TOWN': 'HUDDERSFIELD',
    'LEVANTE UD': 'LEVANTE',
    'MALAGA CF': 'MÁLAGA',
    'MANCHESTER CITY': 'MAN CITY',
    'MANCHESTER UNITED': 'MAN UTD',
    'NEWCASTLE UNITED': 'NEWCASTLE',
    'NORWICH CITY': 'NORWICH',
    'RCD MALLORCA': 'MALLORCA',
```

```

'REAL BETIS BALOMPIE': 'REAL BETIS',
'REAL VALLADOLID CF': 'VALLADOLID',
'SC FREIBURG': 'FREIBURG',
'SC PADERBORN 07': 'PADERBORN',
'SD EIBAR': 'EIBAR',
'SD HUESCA': 'HUESCA',
'SHEFFIELD UNITED': 'SHEFFIELD UTD',
'SPORTING GIJON': 'GIJÓN',
'STOKE CITY': 'STOKE',
'SV DARMSTADT 98': 'DARMSTADT',
'SV WERDER BREMEN': 'W. BREMEN',
'SWANSEA CITY': 'SWANSEA',
'TOTTENHAM HOTSPUR': 'TOTTENHAM',
'TSG 1899 HOFFENHEIM': 'HOFFENHEIM',
'UD LAS PALMAS': 'LAS PALMAS',
'VALENCIA CF': 'VALENCIA',
'VFB STUTTGART': 'STUTTGART',
'VFL WOLFSBURG': 'WOLFSBURG',
'VILLARREAL CF': 'VILLARREAL',
'WEST BROMWICH ALBION': 'WEST BROM',
'WEST HAM UNITED': 'WEST HAM',
'WOLVERHAMPTON WANDERERS': 'WOLVES'
}

```

```

In [24]: # Replace all values in additional set with team names as formatted in original set
for k, v in team_name_key.items():
    df_addition.replace(k, v, inplace=True)

# Verify that all rows now have equivalent team names in the original set
df_addition.loc[~df_addition.Team.isin(df_orig['Home Team'])].empty

```

Out[24]: True

With our renaming successful, we've now accounted for each team in our additional dataset to match with its corresponding format in the original dataset.

Merge Preparation and Verification

To prepare to verify our two datasets are consistent for merging, we'll add calculated columns containing win data to our original dataset. This will be helpful to check season records against match level data.

```

In [25]: # Add booleans for wins, losses, and draws based on the original dataset
df_orig['home_team_win'] = df_orig['Home Team Goals Scored'] > df_orig['Away Team Goals Scored']
df_orig['away_team_win'] = df_orig['Home Team Goals Scored'] < df_orig['Away Team Goals Scored']
df_orig['draw'] = df_orig['Home Team Goals Scored'] == df_orig['Away Team Goals Scored']
df_orig[['Score', 'home_team_win', 'away_team_win', 'draw']].head()

```

```

Out[25]:
   Score  home_team_win  away_team_win  draw
0    1-2             False             True  False
1    2-2             False             False   True
2    2-2             False             False   True
3    0-1             False             True  False
4    0-1             False             True  False

```

To ensure we're merging team data on the correct season, we'll use our calculated wins, losses, and draws to assert our data is consistent between the datasets

In [26]:

```
# Isolate only seasons and teams located in both datasets
years = [year for year in df_orig.year.unique() if str(year) in df_addition.Season.apply(lambda x: str(x).split('/')[0])]
teams = [team for team in df_orig['Home Team'].unique() if team in df_addition.Team.unique()]

# Iterate for each datapoint to check that our match win data is consistent between datasets
for year in years:
    for team in teams:
        # Obtain all matches played by a particular team in a particular season
        orig_team_df = df_orig.loc[((df_orig['Home Team'] == team) | (df_orig['Away Team'] == team)) & (df_orig['Season'] == year)]

        # Determine number of matches played
        orig_num_games = len(orig_team_df)

        # Go ahead and break this iteration if no data is available for this team in this season
        if not orig_num_games:
            continue

        # Determine number of wins based on match data
        orig_num_win = len(orig_team_df.loc[(orig_team_df['Home Team'] == team) & (orig_team_df['Result'] == 'W')])

        # Determine number of losses based on match data
        orig_num_loss = len(orig_team_df.loc[(orig_team_df['Home Team'] == team) & (orig_team_df['Result'] == 'L')])

        # Determine number of draws based on match data
        orig_num_draw = len(orig_team_df.loc[(orig_team_df['Home Team'] == team) & (orig_team_df['Result'] == 'D')])

        # Isolate corresponding row in the additional dataset
        additional_row = df_addition.loc[(df_addition.Team == team) & (df_addition.Season == year)]

        # Test assertions that our calculations are consistent with the additional data
        try:
            assert orig_num_games == additional_row.Matches
            assert orig_num_win == additional_row.Wins
            assert orig_num_loss == additional_row.Losses
            assert orig_num_draw == additional_row.Draws
        except AssertionError:
            print(f'Mismatch for {team} in season {year}')
```

Because the exception block does not execute, we've confirmed that these data are consistent and can continue to merge.

Completing the Merge

To begin, we'll isolate the variables we'd like to merge with our original data set and reformat the season column.

In [27]:

```
# Drop the league column and reformat the season column to match original dataset
df_addition.drop('League', axis=1, inplace=True)
df_addition.Season = df_addition.Season.apply(lambda x: int(x.split('/')[0]))
```

Now, we'll create two copies of the data frame to merge season record data on both home and away teams

In [28]:

```
# Create two copies of the dataframe to contain columns for Home and Away teams
home_season_record = df_addition.copy()
away_season_record = df_addition.copy()
```

In [29]:

```
# Rename the columns of both dataframe copies to be more informative alongside the original
for df, var in [(home_season_record, 'Home'), (away_season_record, 'Away')]:
    cols = {
```

```

'Season': 'year',
'Position': f'{var} Team Season Position',
'Team': f'{var} Team',
'Matches': f'{var} Team Season Matches Played',
'Wins': f'{var} Team Season Wins',
'Draws': f'{var} Team Season Draws',
'Losses': f'{var} Team Season Losses',
'GoalsScored': f'{var} Team Season Goals Scored',
'GoalsConceded': f'{var} Team Season Goals Conceded',
'GoalDiff': f'{var} Team Season Goal Difference',
'FinalPoints': f'{var} Team Season Final Points'
}
df.rename(columns=cols, inplace=True)

```

Finally, we'll create our merged dataset containing season record data for both teams in each match.

In [30]:

```

# Merge data for both Home and Away teams
df_merged = df_orig.merge(home_season_record, on=['year', 'Home Team']).merge(away_season_
df_merged.head()

```

Out[30]:

	Unnamed: 0	Home Team	Away Team	Score	Half Time Score	Match Excitement	Home Team Rating	Away Team Rating	Home Team Possession %	Away Team Possession %
0	380	MAN UTD	TOTTENHAM	1-0	1-0	3.1	6.8	6.4	50	50
1	579	EVERTON	TOTTENHAM	1-1	1-1	5.4	6.7	6.7	41	50
2	611	NORWICH	TOTTENHAM	0-3	0-2	6.4	4.9	7.7	46	50
3	403	LEICESTER CITY	TOTTENHAM	1-1	0-0	5.6	6.4	6.4	35	60
4	478	BOURNEMOUTH	TOTTENHAM	1-5	1-3	6.5	4.6	8.7	46	50

5 rows x 63 columns

In [31]:

```
df_merged.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 5024 entries, 0 to 5023
Data columns (total 63 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Unnamed: 0                               5024 non-null   int64
1   Home Team                                5024 non-null   object
2   Away Team                                5024 non-null   object
3   Score                                    5024 non-null   object
4   Half Time Score                          5024 non-null   object
5   Match Excitement                         5024 non-null   float64
6   Home Team Rating                         5024 non-null   float64
7   Away Team Rating                        5024 non-null   float64
8   Home Team Possession %                   5024 non-null   int64
9   Away Team Possession %                   5024 non-null   int64
10  Home Team Off Target Shots                5024 non-null   float64
11  Home Team On Target Shots                 5024 non-null   float64
12  Home Team Total Shots                     5024 non-null   float64
13  Home Team Blocked Shots                   5024 non-null   float64
14  Home Team Corners                         5024 non-null   float64
15  Home Team Throw Ins                       5024 non-null   float64
16  Home Team Pass Success %                  5024 non-null   float64
17  Home Team Aerials Won                     5024 non-null   float64
18  Home Team Clearances                      5024 non-null   float64

```

```

19 Home Team Fouls 5024 non-null float64
20 Home Team Yellow Cards 5024 non-null float64
21 Home Team Second Yellow Cards 5024 non-null float64
22 Home Team Red Cards 5024 non-null float64
23 Away Team Off Target Shots 5024 non-null float64
24 Away Team On Target Shots 5024 non-null float64
25 Away Team Total Shots 5024 non-null float64
26 Away Team Blocked Shots 5024 non-null float64
27 Away Team Corners 5024 non-null float64
28 Away Team Throw Ins 5024 non-null float64
29 Away Team Pass Success % 5024 non-null float64
30 Away Team Aerials Won 5024 non-null float64
31 Away Team Clearances 5024 non-null float64
32 Away Team Fouls 5024 non-null float64
33 Away Team Yellow Cards 5024 non-null float64
34 Away Team Second Yellow Cards 5024 non-null float64
35 Away Team Red Cards 5024 non-null float64
36 Home Team Goals Scored 5024 non-null int64
37 Away Team Goals Scored 5024 non-null int64
38 Home Team Goals Conceded 5024 non-null int64
39 Away Team Goals Conceded 5024 non-null int64
40 year 5024 non-null int64
41 league 5024 non-null object
42 home_team_win 5024 non-null bool
43 away_team_win 5024 non-null bool
44 draw 5024 non-null bool
45 Home Team Season Position 5024 non-null object
46 Home Team Season Matches Played 5024 non-null int64
47 Home Team Season Wins 5024 non-null int64
48 Home Team Season Draws 5024 non-null int64
49 Home Team Season Losses 5024 non-null int64
50 Home Team Season Goals Scored 5024 non-null int64
51 Home Team Season Goals Conceded 5024 non-null int64
52 Home Team Season Goal Difference 5024 non-null int64
53 Home Team Season Final Points 5024 non-null int64
54 Away Team Season Position 5024 non-null object
55 Away Team Season Matches Played 5024 non-null int64
56 Away Team Season Wins 5024 non-null int64
57 Away Team Season Draws 5024 non-null int64
58 Away Team Season Losses 5024 non-null int64
59 Away Team Season Goals Scored 5024 non-null int64
60 Away Team Season Goals Conceded 5024 non-null int64
61 Away Team Season Goal Difference 5024 non-null int64
62 Away Team Season Final Points 5024 non-null int64
dtypes: bool(3), float64(29), int64(24), object(7)
memory usage: 2.4+ MB

```

Now, we've accumulated a single dataframe containing both match data alongside season records for both home and away teams.

Part 5 - Explore Bivariate relationships

- Choose a reasoned set of variables to explore further. You don't have to explore all possible pairs of variables, nor do we want to grade that much. Choose 7 - 9 variables. One should be a variable that you'd like to predict (target variable) using the others (predictor variables).
- List your predictor variables
- List your target variable
- Briefly describe why you have chosen these.

Use any of the available visualizations from Seaborn to explore the relationships between the variables. Explore the relationships among the predictor variables as well as the relationship between each predictor variable and the target variable. Which of the predictor variables are most strongly related? Are there any

interesting relationships between categorical predictors and numeric predictors? If there are any dichotomous variables, does that influence any of the relationships? Are the relationships positive or negative?

Below each plot, you should provide a description and interpretation of the plot. Make sure to include why the variables in that plot were chosen and what you hope the reader would gain from it as well.

Relationship 1: Average difference in possession as a predictor for season ranking

Predictor variable: $\frac{\sum_{j=1}^k (2P_j - 100)}{k}$ where P_j represents a team's possession percentage on $[0, 100]$ percent of the j th game of k games in a given season.

Target variable: $1 - \frac{R}{L}$ where R represents a team's season rank out of L total ranked teams in their league.

Explanation: With this relationship, we wish to explore whether a team's tendency to play more offensively or defensively (positive and negative average difference of possession, respectively) influences the team's final position of the league's rankings.

In [32]:

```
# Obtain individual team match statistics agnostic of home vs. away
df_1 = df_merged[['year', 'league', 'Home Team', 'Home Team Possession %', 'Home Team Season Rank']]
df_2 = df_merged[['year', 'league', 'Away Team', 'Away Team Possession %', 'Away Team Season Rank']]
for df, var in [(df_1, 'Home'), (df_2, 'Away')]:
    df.rename(
        columns={
            f'{var} Team': 'team',
            f'{var} Team Possession %': 'pos_diff',
            f'{var} Team Season Rank': 'rank'
        },
        inplace=True
    )

# Concatenate the team statistics together and calculate difference in ball possession
df_comb = pd.concat([df_1, df_2])
df_comb.rank = df_comb.rank.apply(lambda x: int(x))
df_comb['pos_diff'] = 2 * df_comb.pos_diff - 100

# Group by team and season to obtain average possession differences
df_comb = df_comb.drop('pos_diff', axis=1).groupby(['year', 'league', 'team']).mean()

# Calculate rank percentile
max_rank = df_comb[['year', 'league', 'rank']].groupby(['year', 'league']).max()
df_comb['pctile'] = pd.Series([1 - v.rank / max_rank.loc[v.year, v.league].rank for v in df_comb.itertuples()])
df_comb.head()
```

Out[32]:

	year	league	team	season_rank	pos_diff	pctile
0	2015	english	ARSENAL	2.0	16.631579	0.90
1	2015	english	ASTON VILLA	20.0	-7.263158	0.00
2	2015	english	BOURNEMOUTH	16.0	2.421053	0.20
3	2015	english	CHELSEA	10.0	12.315789	0.50
4	2015	english	CRYSTAL PALACE	15.0	-8.421053	0.25

In [33]:

```
# Plot average possession difference vs. overall season rank percentile
```



```
sns.scatterplot(x=df_comb.pos_diff, y=df_comb.pctile)
pass
```



This plot demonstrates the relationship between a team's gameplay style on the x-axis (with negative values being mostly defensive during a season and positive values being mostly offensive during a season), and a team's percentile rank at the end of the season.

From the plot, we see a fairly strong positive correlation indicating that teams who control the ball for more gameplay during a season have a better likelihood of performing well in their league. Conversely, teams who play mostly defensively during a season rank poorly among other teams in their league.

Relationship 2: Season Performance for Relative Match Excitement

Predictor variable: End-of-Season Ranking

Target variable: Match Excitement

Explanation: With this relationship, we wish to explore whether the season rank of a team generates more match excitement.

In [34]:

```
# Obtain individual team match statistics agnostic of home vs. away
df_3 = df_merged[['year', 'Home Team', 'Match Excitement', 'Home Team Season Position']]
df_4 = df_merged[['year', 'Away Team', 'Match Excitement', 'Away Team Season Position']]
for df, var in [(df_3, 'Home'), (df_4, 'Away')]:
    df.rename(
        columns={
            f'{var} Team': 'team',
            f'{var} Team Season Position': 'season_rank'
        },
        inplace=True
    )

# Concatenate the team statistics together
df_comb = pd.concat([df_3, df_4])

# Rename match excitement and convert season rank to int datatype (used to be as string)
df_comb.rename(columns={'Match Excitement': 'match_excitement'}, inplace=True)
```

```
df_comb['season_rank'] = df_comb['season_rank'].astype(int)
df_comb

df_comb.groupby(['team', 'year']).mean().reset_index().drop(['team', 'year'], axis=1)


# Sort the dataframe by the season rank
df_comb = df_comb.sort_values(by='season_rank')
df_comb.head()
```

Out [34]:

	year	team	match_excitement	season_rank
2359	2017	BAYERN	8.7	1
3517	2016	REAL MADRID	9.6	1
4080	2017	BARCELONA	7.5	1
4465	2018	BARCELONA	4.3	1
3312	2015	BARCELONA	3.8	1

In [35]:

```
# Create a scatterplot showing the relation between season rank and match excitement
sns.scatterplot(data=df_comb, x="season_rank", y="match_excitement", alpha=0.2, s=150)
pass
```



After analyzing this scatterplot, we came to the conclusion that there is no direct correlation between the ranking of a team and the excitement generated during a match. This could be due to a variety of reasons, such as the crowd rooting for underdogs or players being transferred to lower ranking teams. Overall, this proves that our hypothesis related to match excitement being generated by ranking was incorrect.

Relationship 3: Fouls, Yellow, and Red Cards Collected for End-of-Season Rank

Predictor variable: The "Aggressive Score" which is an aggregate score of fouls, yellow cards, second yellow cards, and red cards, each representing a differing magnitude of offense.

Target variable: The End-of-Season Rank of a team within their respective league.

Explanation: With this relationship, we wish to explore whether the total quantity of yellow, second-yellow, red cards, and fouls collected throughout a season affects a team's final rank.

The driving question as we explore this relationship is whether having aggressive players on a team drives season performance.

In [36]:

```
# Obtain individual team match statistics agnostic of home vs. away
df_5 = df_merged[['Home Team', 'Home Team Yellow Cards', 'Home Team Second Yellow Cards',
df_6 = df_merged[['Away Team', 'Away Team Yellow Cards', 'Away Team Second Yellow Cards',

for df, var in [(df_5, 'Home'), (df_6, 'Away')]:
    df.rename(
        columns={
            f'{var} Team': 'team',
            f'{var} Team Season Position': 'season_rank',
            f'{var} Team Yellow Cards': 'yellow_cards',
            f'{var} Team Second Yellow Cards': 'second_yellow_cards',
            f'{var} Team Red Cards': 'red_cards',
            f'{var} Team Fouls': 'fouls',
        },
        inplace=True
    )

# Concatenate the team statistics together
df_comb = pd.concat([df_5, df_6])
df_comb['season_rank'] = df_comb['season_rank'].astype(int)

df_comb['aggressive_score'] = df_comb['fouls'] + 2 * df_comb['yellow_cards'] + 3 * df_comb

df_comb.sort_values(by='aggressive_score', ascending=False)
```

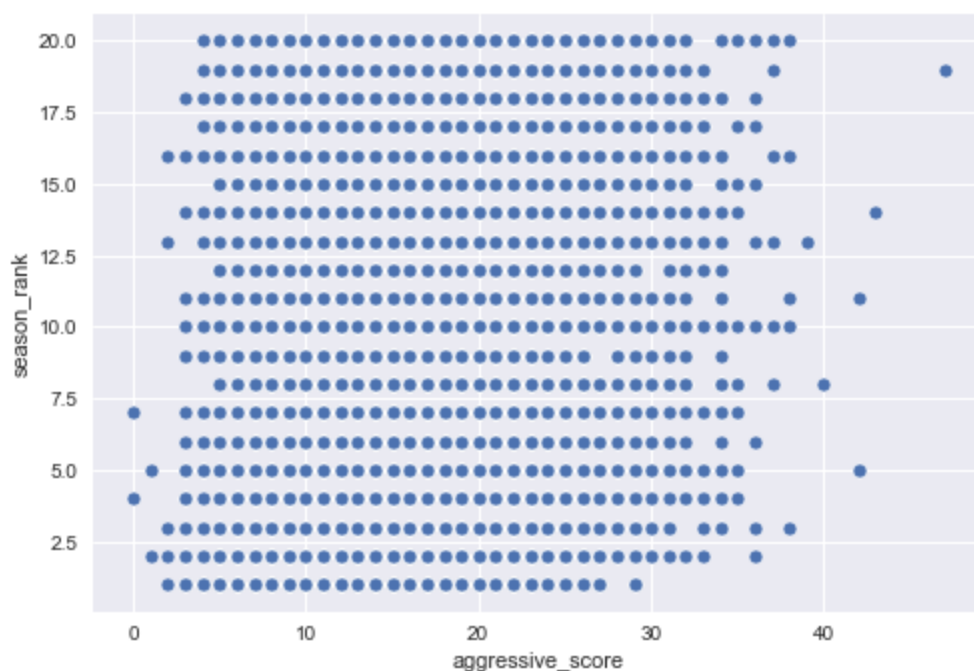
Out [36]:

	team	yellow_cards	second_yellow_cards	red_cards	fouls	season_rank	aggressive_score
3322	GETAFE	8.0	0.0	1.0	27.0	19	47.0
4067	ALAVÉS	7.0	0.0	0.0	29.0	14	43.0
1986	FRANKFURT	5.0	0.0	1.0	28.0	11	42.0
4423	GETAFE	7.0	0.0	0.0	28.0	5	42.0
4428	GETAFE	5.0	2.0	0.0	26.0	5	42.0
...
947	TOTTENHAM	0.0	0.0	0.0	2.0	3	2.0
4486	ATLETICO MADRID	0.0	0.0	0.0	1.0	2	1.0
1589	LEICESTER CITY	0.0	0.0	0.0	1.0	5	1.0
1784	CHELSEA	0.0	0.0	0.0	0.0	4	0.0
1336	WOLVES	0.0	0.0	0.0	0.0	7	0.0

10048 rows x 7 columns

In [37]:

```
# Create scatterplot showing the relationship between how the aggressiveness of a team af
sns.scatterplot(data=df_comb, x="aggressive_score", y="season_rank")
pass
```



From this scatterplot, we can see that there is no clear pattern between aggressiveness of a team and their final season ranking. However, notably, some higher ranked teams tend to have a rather high "Aggressive Score". Since a lot of the lower ranking teams are centralized in the 15-25 Aggressive Score range, it appears that teams outside of this ranges perform either have high or low performance.

Relationship 4: A team's rating as a predictor for match win.

Predictor variable: The difference in two teams' ratings as determined by previous gameplay.

Target variable: The result of the game—either win, lose, or draw.

Explanation: With this relationship, we seek to determine if a team's previous performance as summarized by their rating is indicative of their likelihood of winning a match.

In [38]:

```
# Obtain relevant columns
df_1 = df_merged[['Home Team Rating', 'Away Team Rating', 'home_team_win', 'away_team_win']
df_2 = df_merged[['Away Team Rating', 'Home Team Rating', 'away_team_win', 'home_team_win']]

# Reformat column names and concatenate into single dataframe
renamed_cols = ['rating', 'op_rating', 'win', 'lose', 'draw']
df_1.columns = renamed_cols
df_2.columns = renamed_cols
df_comb = pd.concat([df_1, df_2])

# Calculate difference in rating vs. opponent rating
df_comb['rating_diff'] = df_comb.rating - df_comb.op_rating
df_comb.head()
```


Out [38]:

	rating	op_rating	win	lose	draw	rating_diff
0	6.8	6.4	True	False	False	0.4
1	6.7	6.7	False	False	True	0.0
2	4.9	7.7	False	True	False	-2.8
3	6.4	6.4	False	False	True	0.0
4	4.6	8.7	False	True	False	-4.1

In [39]:

```
# Setup multiplot
fig, axs = plt.subplots(ncols=3)
axs[0].set_ylabel('Difference in Team Rating Prior to Match')
for ax in axs:
    ax.set_ylim(-9, 9)

# Plot a violin plot for wins, loses, and draws
for i, (col, color, label) in enumerate([
    (df_comb.win, 'forestgreen', 'Win'),
    (df_comb.lose, 'firebrick', 'Lose'),
    (df_comb.draw, 'goldenrod', 'Draw')
]):
    sns.violinplot(data=df_comb[['rating_diff']].loc[col], scale='count', ax=axs[i], color=color)
    axs[i].set_xlabel(label)
```



In these plots, we gain a sense of how a team's prior performance is a good indicator of their odds of winning a game. Positive differences (teams that are higher rated than their opponents) tend to win, with upset occurrences being uncommon. Interestingly, draws will most frequently occur between teams that are "evenly matched," as far as the ratings are calculated.

Relationship 5: Team ahead at halftime indicating winner

Predictor variable: The team winning at halftime

Target variable: The team winning the match

Explanation: For this relationship, we'll see how often the results at halftime carry over to the results following the match

In [40]:

```
# Obtain relevant columns and augment dataframe to determine frequency of match outcomes
df = df_merged[['Half Time Score', 'Score']]
df.columns = ['at_half', 'final']

get_winner = lambda x: 'HOME' if x.split('-')[0] > x.split('-')[1] else 'AWAY' if x.split(
df = df.apply(lambda x: x.apply(get_winner)).pivot_table(index='at_half', columns='final',
df
```

Out[40]:

	final	AWAY	DRAW	HOME
at_half				
AWAY	889	221	127	
DRAW	500	756	793	
HOME	98	256	1384	


In [41]:

```
# Setup polar grid for radar plot
angles = np.linspace(0, 2 * np.pi, 3, endpoint=False).tolist()
angles += angles[:1]
fig, ax = plt.subplots(figsize=(6, 6), subplot_kw=dict(polar=True))

# Plot each match outcome
for col, color, label in [
    ('HOME', 'mediumpurple', 'Home Ahead at Half'),
    ('AWAY', 'deepskyblue', 'Away Ahead at Half'),
    ('DRAW', 'darksalmon', 'Game Tied at Half')
]:
    values = df.loc[col].tolist()
    values += values[:1]
    ax.plot(angles, values, color=color, linewidth=1, label=label)
    ax.fill(angles, values, color=color, alpha=0.2)

# Rotate plot and add labels
ax.set_theta_offset(np.pi / 2)
ax.set_theta_direction(-1)
ax.set_thetagrids(np.degrees(angles), [None, 'Match Ends in Draw', 'Home Team Wins', 'Away
ax.set_rlabel_position(60)
ax.legend(loc='lower center')
pass

# Radar plot reference: https://www.pythoncharts.com/matplotlib/radar-charts/
```



From this radar plot, we see that a team ahead at the half will usually go on to win the match. Games tied at the half are largely a tossup. Additionally, from the magnitude of each polygon's extreme point, we note a home team advantage. In this regard, a home team is more likely to win the match despite losing at the half.